

MySQL与CPU

@淘宝丁奇



追風堂



讲师简介：

丁奇：

08年至10年在百度贴吧，作服务端开发，开始接触MySQL。之后由于业务需要开始看MySQL代码，囫圇吞枣不求甚解。10年得机会进入淘宝核心系统数据库组，主要是MySQL优化和提升可维护性。参与IC、TC读库调优；写了一些插件，打了几个patch到官方；实现MySQL主从同步工具、设计MySQL异构数据同步方案、MySQL中间层。一直游离在了解需求、设计方案、推广方案的三点一线上。

课程目标与目标学员页



- **目标学员：**对MySQL、系统优化有兴趣的同学。
- **课程目的：**介绍MySQL的线程状况和cpu
- **学员能够获得的收获：**了解MySQL的各个线程工作流程、CPU压力的应用下的瓶颈分析、一些追查问题的方法。






1. 线程与CPU

- MySQL里有哪些线程，各司何职？
- 与线程有关的参数
- 查询、更新涉及到哪些线程
- 并发线程和性能的关系

2. 单线程内的CPU消耗

- 如何分析CPU消耗
- 排序操作与CPU，优化空间
- 大小头问题与CPU，优化方法
- Spinlock问题




MySQL里有哪些线程，各司何职？

1. 怎么看程序线程？

```
Thread 24 (Thread 0x415b3940 (LWP 26605)):  
#0  0x00000033d660ab99 in pthread_cond_wait@@GLIBC_2.3.2 ()  
#1  0x000000000089f1b7 in os_event_wait_low ()  
#2  0x000000000089dd3e in os_aio_simulated_handle ()  
#3  0x000000000085ed96 in fil_aio_wait ()  
#4  0x00000000007f3d48 in io_handler_thread ()  
#5  0x00000033d66064a7 in start_thread () from /lib64/libpthread.so.0  
#6  0x00000033d5ed3c2d in clone () from /lib64/libc.so.6  
Thread 23 (Thread 0x4223a940 (LWP 26606)):  
#0  0x00000033d660ab99 in pthread_cond_wait@@GLIBC_2.3.2 ()  
#1  0x000000000089f1b7 in os_event_wait_low ()  
#2  0x000000000089dd3e in os_aio_simulated_handle ()  
#3  0x000000000085ed96 in fil_aio_wait ()  
#4  0x00000000007f3d48 in io_handler_thread ()  
#5  0x00000033d66064a7 in start_thread () from /lib64/libpthread.so.0  
#6  0x00000033d5ed3c2d in clone () from /lib64/libc.so.6  
Thread 22 (Thread 0x42c3b940 (LWP 26607)):  
#0  0x00000033d660ab99 in pthread_cond_wait@@GLIBC_2.3.2 ()  
#1  0x000000000089f1b7 in os_event_wait_low ()
```

MySQL里有哪些线程，各司何职？

1. 当出现死锁等情况时，也可以用pstack各个线程在忙什么。
2. 启动后的18个线程，分别是做什么的？
3. 来一个新客户端连接什么情况？
4. 断开一个客户端连接什么情况？--为什么线程没少？



初始线程分析

io_handler_thread 10个

Thread 18 (Thread 0x42816940 (LWP 22173)):

#0 0x00000033d660ab99 in pthread_cond_wait@@GLIBC_2.3.2
#1 0x0000000000008c5797 in os_event_wait_low ()
#2 0x0000000000008c430e in os_aio_simulated_handle ()
#3 0x00000000000087f89e in fil_aio_wait ()
#4 0x000000000000806638 in io_handler_thread ()
#5 0x000000033d66064a7 in start_thread () from /lib64/libc.so.6
#6 0x000000033d5ed3c2d in clone () from /lib64/libc.so.6



初始线程分析

io_handler_thread 10个

```
/* Now overwrite the value on srv_n_file_io_threads */  
srv_n_file_io_threads = 2 + srv_n_read_io_threads  
                      + srv_n_write_io_threads;
```

```
/* Create i/o-handler threads: */  
  
for (i = 0; i < srv_n_file_io_threads; i++) {  
    n[i] = i;  
  
    os_thread_create(io_handler_thread, n + i, thread_ids + i);  
}
```

```
while (srv_shutdown_state != SRV_SHUTDOWN_EXIT_THREADS) {  
    fil_aio_wait(segment);  
}
```

Insert_buffer \ log \ read\write thread



初始线程分析

Srv.*thread 6个

1. srv_lock_timeout_thread
2. srv_error_monitor_thread
3. srv_monitor_thread
4. srv_LRU_dump_restore_thread (Percona)
5. srv_master_thread
6. srv_purge_thread (5.5)



初始线程分析

剩下两个

所以初始时MySQL层就两个线程

1. main 线程
2. signal_handler

新连接

从库谁管？

Thread 2 (Thread 0x41c7e940 (LWP 26056)):

#0 0x00000033d660ab99 in pthread_cond_wait@

#1 0x000000000004fbd4a in one_thread_per_connection_end ()

#2 0x000000000005f85ef in do_handle_one_connection ()

#3 0x000000000005f871a in handle_one_connection ()

#4 0x00000033d66064a7 in start_thread () from

/lib64/libpthread.so.0

#5 0x00000033d5ed3c2d in clone () from

两个例子

```
Thread 149 (Thread 0x4b5be940 (LWP 1075)):  
#0  0x0000003d1c40ab99 in pthread_cond_wait@@GLIBC_2.3.2 ()  
#1  0x0000000000537b62 in THD::wait_for_wakeup_ready ()  
#2  0x00000000006e7c8a in MYSQL_BIN_LOG::write_transaction_to_binlog_events ()  
#3  0x00000000006e7efc in MYSQL_BIN_LOG::write_transaction_to_binlog ()  
#4  0x00000000006e8727 in MYSQL_BIN_LOG::log_and_order ()  
#5  0x000000000065a9c3 in ha_commit_trans ()  
#6  0x00000000005fbbc2 in trans_commit_stmt ()  
#7  0x00000000006f05aa in rows_event_stmt_cleanup ()  
#8  0x00000000006f857c in Rows_log_event::do_apply_event ()  
#9  0x00000000005dae28 in mysql_client_binlog_statement ()  
#10 0x000000000055a412 in mysql_execute_command ()  
#11 0x000000000055e4ca in mysql_parse ()  
#12 0x000000000055f82b in dispatch_command ()  
#13 0x000000000055fb6e in do_command ()  
#14 0x00000000005ef2a2 in do_handle_one_connection ()  
#15 0x00000000005ef37a in handle_one_connection ()  
#16 0x0000003d1c4064a7 in start_thread () from /lib64/libpthread.so.0  
#17 0x0000003d1b8d3c2d in clone () from /lib64/libc.so.6
```




两个例子

```
#1 0x00002aaac4d7b7c5 in os_event_wait_low () from /opt/soft/mysql/lib/mysql/plugin/ha_inn
#2 0x00002aaac4dc26ae in sync_array_wait_event () from /opt/soft/mysql/lib/mysql/plugin/ha
#3 0x00002aaac4dc3796 in mutex_spin_wait () from /opt/soft/mysql/lib/mysql/plugin/ha_innod
#4 0x00002aaac4dd53c8 in trx_commit_for_mysql () from /opt/soft/mysql/lib/mysql/plugin/ha
#5 0x00002aaac4da6a6c in row_drop_table_for_mysql () from /opt/soft/mysql/lib/mysql/plugin
#6 0x00002aaac4d531b7 in ha_innodb::delete_table(char const*) () from /opt/soft/mysql/lib/
#7 0x000000000006d2759 in ha_delete_table(THD*, handlerton*, char const*, char const*, char
#8 0x000000000006e05de in quick_rm_table(handlerton*, char const*, char const*, unsigned in
#9 0x000000000006e9ef9 in mysql_alter_table(THD*, char*, char*, st_ha_create_information*,
#10 0x000000000005e7f6e in mysql_execute_command(THD*) ()
#11 0x0000000000073696c in sp_instr_stmt::exec_core(THD*, unsigned int*) ()
#12 0x00000000000739347 in sp_lex_keeper::reset_lex_and_exec_core(THD*, unsigned int*, bool,
#13 0x0000000000073ba08 in sp_instr_stmt::execute(THD*, unsigned int*) ()
#14 0x0000000000073a8c1 in sp_head::execute(THD*) ()
#15 0x0000000000073c653 in sp_head::execute_procedure(THD*, List<Item>*) ()
#16 0x000000000005e7ac1 in mysql_execute_command(THD*) ()
#17 0x0000000000073696c in sp_instr_stmt::exec_core(THD*, unsigned int*) ()
#18 0x00000000000739347 in sp_lex_keeper::reset_lex_and_exec_core(THD*, unsigned int*, bool,
#19 0x0000000000073ba08 in sp_instr_stmt::execute(THD*, unsigned int*) ()
#20 0x0000000000073a8c1 in sp_head::execute(THD*) ()
#21 0x0000000000073c653 in sp_head::execute_procedure(THD*, List<Item>*) ()
#22 0x0000000000074aca8 in Event_job_data::execute(THD*, bool) ()
#23 0x00000000000749303 in Event_worker_thread::run(THD*, Event_queue_element_for_exec*) ()
#24 0x00000000000749358 in event_worker_thread ()
#25 0x00000003a8b60673d in start_thread () from /lib64/libpthread.so.0
#26 0x00000003a8aed44bd in clone () from /lib64/libc.so.6
```

MySQL里有哪些线程，各司何职？

1. 当出现死锁等情况时，也可以用pstack各个线程在忙什么。
2. 启动后的18个线程，分别是做什么的？
3. 来一个新客户端连接什么情况？
4. 断开一个客户端连接什么情况？--为什么线程没少？




与线程数目有关的参数

- innodb_purge_threads
- innodb_read_io_threads / innodb_write_io_threads

与线程控制有关的参数

- thread_cache_size
- thread_concurrency （已经废弃）
- innodb_thread_concurrency
- thread_handling (one-thread-per-connection)



查询涉及到哪些线程




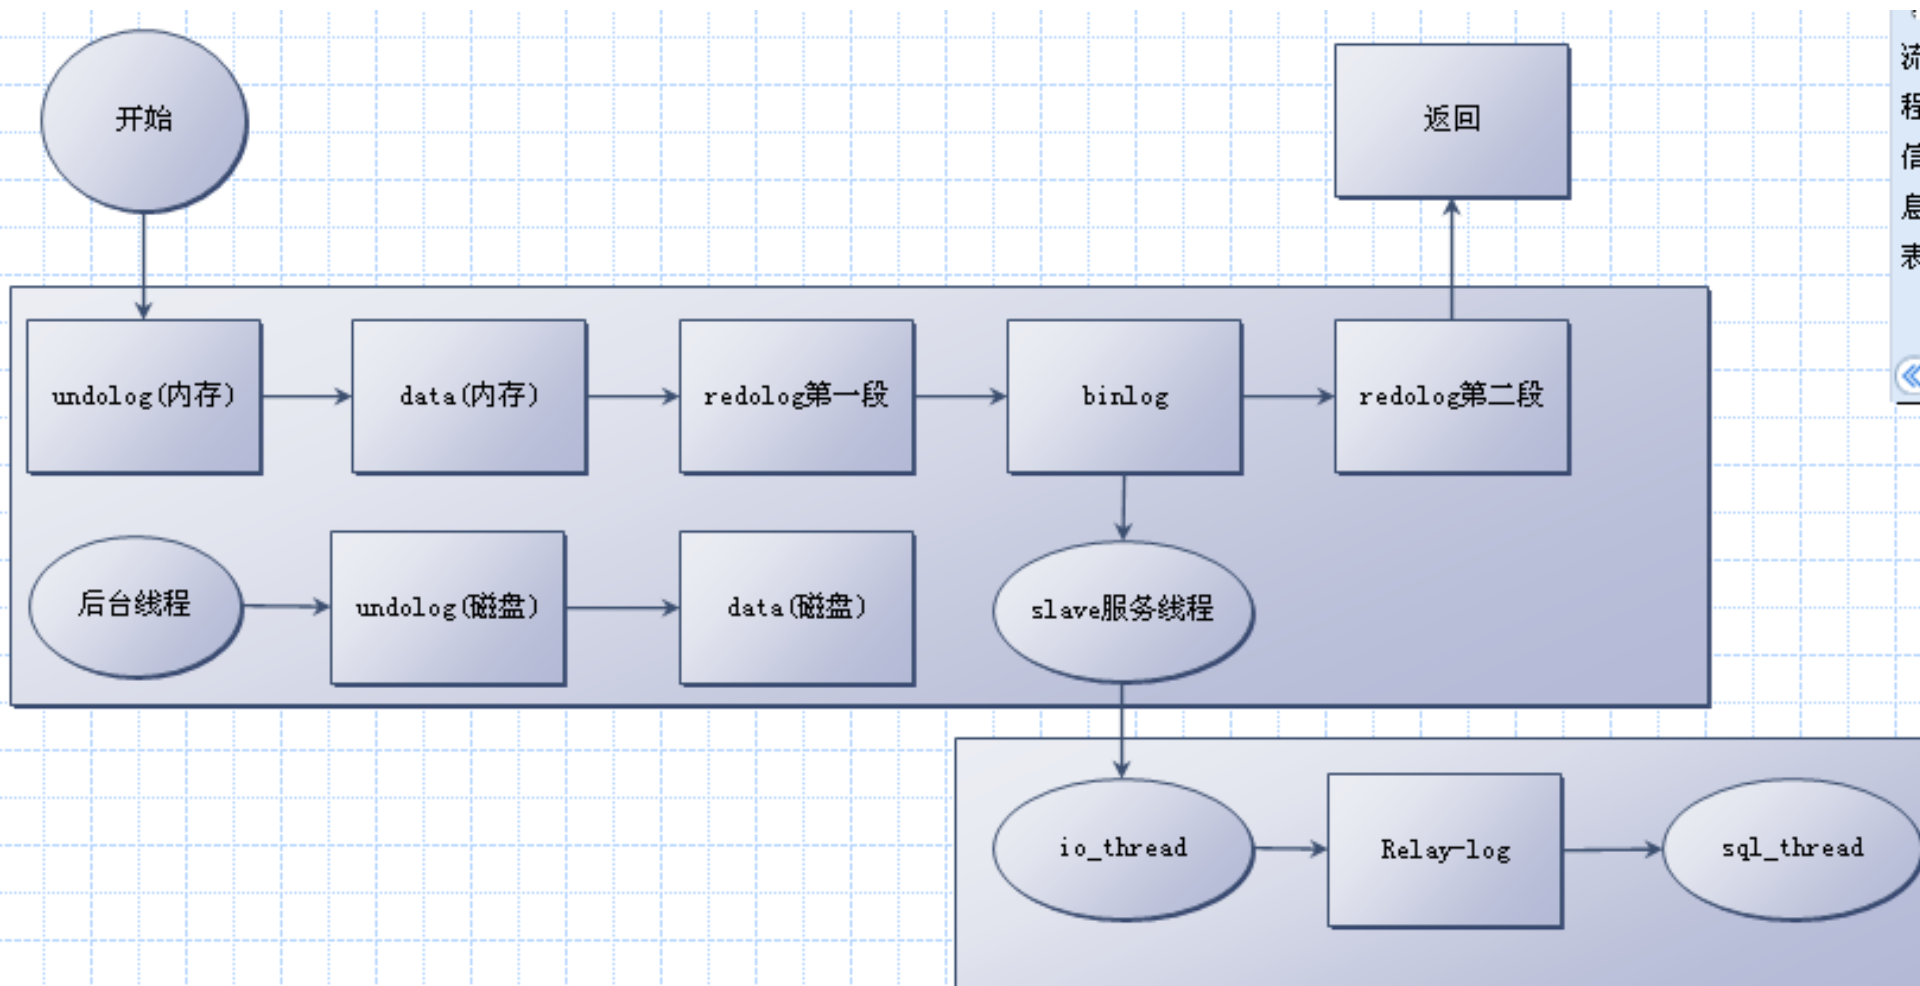
访问数据的方式与此有关

- Select 1 行
- Select 全表

异步read的触发时机 / innodb_read_ahead_threshold



更新涉及到哪些线程



更新涉及到哪些线程



除用户线程外

- Insert buffer thread
- Log thread
- Io write thread
- Master thread
- Purge thread



并发线程和性能的关系



要达到最大性能，单线程好还是多线程？

- 多核机器当然要多线程
- 但是不是越多越好
- 线上一个例子及解决方法（治标 or 治本）



查询的各个cpu消耗阶段 及 一些优化手段



➤ SQL 解析


- 注意mysql_parse 不只是解析
- Batch 操作

➤ 数据查询

- querycache
- 索引、排序
- Spin lock

➤ 数据拷贝

- 引擎层与Server层拷贝，大小头问题



大小头问题



- 现象
- 导致的cpu问题
- 改进空间

Spin lock

- 什么是spinlock
- 工作过程
- 什么情况下会成为cpu杀手
- 解决方向



排序操作与CPU，优化空间



- 什么是排序
- 有哪几种排序
- 执行过程
- Cpu都耗在哪里？
- Cpu消耗的模式下可以怎么改进



Batch操作对性能的影响



- 像MySQLdump那样的insert
- Load data的实际流程
- Transfer增加batchinsert的思路和效果
- Analyse table 的流程和优化—fast index creation?



谢谢



追風堂

